

The prompt is the protocol: a disclosure standard for clinician-investigators using agentic large language models

Hsieh-Ting Lin, M.D.^{1,*}

¹Department of Hematology and Medical Oncology, Koo Foundation Sun Yat-Sen Cancer Center, Taipei, Taiwan. ORCID: 0009-0002-3974-4528.

*Correspondence: mail@hsiehting.com.

Article type. Perspective. **Body word count.** 620 (ceiling 1 200; brevity is consistent with the journal’s “brief, accessible” Perspective format). **References.** 5/5. **Figures.** 1/1. **Conflicts of interest.** None. **Funding.** None.

Abstract. A single clinician-investigator can in 2026 assemble an end-to-end agentic-LLM research workflow on commodity tooling; current biomedical-journal AI policies classify most of what such a workflow does as prohibited, and we propose a Disclosure 2.0 standard — prompts, model identifier, commit history, reviewer-subagent transcripts, audit log, tagged release — under which the present Perspective and its supporting null-result case study are themselves produced.

1 The 2026 landscape

In 2026, two peer-reviewed *Nature* papers placed autonomous agentic-LLM research at the centre of scientific publishing. Google DeepMind’s Co-Scientist¹ generated, critiqued and refined biomedical hypotheses, with acute-myeloid-leukaemia drug-repurposing candidates validated *in vitro*. Sakana AI’s end-to-end pipeline² produced a manuscript that passed first-round workshop peer review under blind review; the authors’ Limitations enumerate hallucinated citations and naive ideas. In parallel, a clinician with no large-lab affiliation can in 2026 assemble a workflow that performs every Co-Scientist task on commodity tooling — one domain prompt, a handful of orchestration prompts, an open-source LLM CLI — and produce a manuscript matched to a target journal’s specifications. This Perspective and its supporting hepatocellular-carcinoma case study are the existence proof.

2 The disclosure gap

The biomedical-journal AI policy in current public form across the Lancet group, NEJM family and JAMA Network permits “readability and language improvements only”³ and prohibits using AI to generate scientific arguments, draft methodology, write literature reviews, or create new content. Read literally, the policy classifies almost every action a 2026 agentic workflow performs as prohibited. Read sympathetically, it protects against three real harms — authorship inflation, scientific hallucination and undocumented intervention — with rules calibrated for a 2023 generation of LLM use. The policy’s intent is sound; the mechanism is mis-calibrated. Disclosure by a single-sentence acknowledgement cannot tell a reader which sentences were AI-drafted, which prompts produced them, or whether the cited literature actually exists.

3 Disclosure 2.0

We propose that the published unit for AI-assisted biomedical research become a tuple of six artefacts at a tagged release: (1) the prompts that orchestrated the work, verbatim, in order, closed for edits; (2) the model identifier with snapshot date and tool envelope; (3) the git commit history with no

force-pushes or rebases; (4) the reviewer-subagent transcripts per round; (5) the audit log from a sealed re-execution-plus-citation-veracity-plus-statistic-rederivation subagent; (6) the tagged release with a persistent DOI. A minimum viable manifest (items 1, 2, 3, 6) covers the consequential audit surfaces at near-zero infrastructure cost. A vigorous human-in-the-loop school exemplified by the Academic Research Skills toolkit⁴ argues for AI under operator supervision rather than autonomous use; the two regimes are different equilibria with different cost-of-error profiles, and Disclosure 2.0 is the disclosure layer both should share. Figure 1 is the artefact ledger of the present Perspective.

4 The supporting null

The supporting case study runs Disclosure 2.0 on a hepatocellular-carcinoma overall-survival stratification problem outside the operator’s clinical sub-specialty. Layer 1, a single autonomous Claude Code session, produced a 50-gene risk score with development-cohort paired-optimism ΔC -index = 0.063 (95 % CI [0.003, 0.113], $n = 366$). Layer 2 (a sealed audit subagent) filed twelve substantive findings including one blocker (a required annotation file not downloaded by the data-prep script) and two high-severity (an abstract estimator that disagrees with the Methods specification; a stratified-vs-stage-adjusted HR discrepancy). Layer 3 (the operator’s preregistered⁵ external validation against held-out GSE14520 and GSE76427, pooled $n = 333$, 107 events) returned $\Delta C = 0.006$, 95 % CI [−0.008, 0.030] — primary hypothesis *not rejected*; calibration drifts from 1.1 at 1 y to 1.9 at 5 y; IDI ≈ 0 . The development-cohort positive does not transport. The null is the headline; the manifest preserves it together with the twelve audit findings as evidence the regime makes its own failure modes visible.

5 Recommendation

NEJM AI and the wider biomedical-journal community should pilot Disclosure 2.0 as an alternative submission track for AI-deep manuscripts, with a JSON-schema manifest (committed at the repository root and validated at submission via PDF attachment plus linked DOI) and a 90-day trial in which five to ten submissions per quarter attach the manifest. The Perspective the reviewer is reading was assembled under exactly that standard; every artefact the manifest would require is publicly archived at <https://github.com/htlin222/end-to-end> under MIT licence at tag `viewpoint-nejmai-v1.0.0`, with a Zenodo DOI minted on tag push. The medium is the message; whatever the editorial decision, the artefact stands.

Declarations

Conflicts of interest. The author declares no competing interests. **Funding.** None. **Role of the funding source.** No funder; full repository access and final submission decision rest with the author. **AI usage.** Manuscript drafted with Anthropic Claude (Opus 4.7, 1M-context variant) via the Claude Code CLI; per-step disclosure log at `docs/ai-usage-disclosure.md` in the linked repository. No AI tool is listed as an author; no AI tool generated figures. **Data and code availability.** All artefacts at <https://github.com/htlin222/end-to-end> (MIT licence; tag `viewpoint-nejmai-v1.0.0`).

References

- [1] Gottweis J, Weng WH, Daryin A, Tu T, Sirkovic P, Myaskovsky A, et al. Accelerating scientific discovery with Co-Scientist. *Nature*. 2026. PMID 42156544.
- [2] Lu C, Lu C, Lange RT, Yamada Y, Hu S, Foerster J, et al. Towards end-to-end automation of AI research. *Nature*. 2026;651(8107):914-9. PMID 41882133.

- [3] The Lancet group. Editorial policies: use of generative AI and large language models; 2024. Retrieved 2026-05-20. Stands in for the biomedical-journal AI-policy class; NEJM AI, JAMA Network, and Nature portfolio carry materially similar wording. <https://www.thelancet.com/editorial-policies>.
- [4] Imbad, contributors. Academic Research Skills for Claude Code: research, write, review, revise, finalize; 2026. Mature contemporaneous example of the human-in-the-loop agentic-research school cited in Section 3. GitHub release v3.9.2, <https://github.com/Imbad0202/academic-research-skills>.
- [5] Lin HT. Preregistration of Layer-3 external validation, end-to-end repository; 2026. Operator-confirmed SHA; repository tag `viewpoint-nejmai-v1.0.0`. GitHub commit `88d6d15`, file `docs/prereg.md`.

